

Seneca



CVI620/ DPS920

Introduction to Computer Vision

Object Recognition

Seneca College

Vida Movahedi

Overview


- Object Detection methods
 - Keypoint matching
 - Template matching
 - Chamfer matching

Keypoint Matching

3 main steps [2]

1. Search an image and find all keypoints
 - Different methods
 2. Create a descriptor for every keypoint found
 3. Find matches in another image or set of images
 - By comparing the descriptors of keypoints
 - Using a matching measure
- Example:
 - Tracking: matching with the next frame
 - Recognition/ Identification: Match with a set of labelled images/ templates





Step 1: Finding keypoints

- **cv::KeyPoint** – the abstract class for finding (the location of) keypoints

```
class cv::KeyPoint {  
  
    public:  
  
        cv::Point2f pt;          // coordinates of the keypoint  
        float      size;        // diameter of the meaningful keypoint neighborhood  
        float      angle;       // computed orientation of the keypoint (-1 if none)  
        float      response;    // response for which the keypoints was selected  
        int        octave;      // octave (pyramid layer) keypoint was extracted from  
        int        class_id;    // object id, can be used to cluster keypoints by object
```

Step 2: Creating a descriptor

- **cv::Feature2D**
 - The abstract class for finding/ detecting/ computing descriptors
 - Often step 1 & 2 are done simultaneously
 - A descriptor list (A vector)
- Derived classes:
 - Harris detector, Shi-Tomasi detector (GoodFeaturesToTrack)
 - SimpleBlobDetector
 - Also implementations available for FAST, SIFT, SURF, ORB, BRISK, etc.
- Detect/ compute key points:
 - `detect(image, keypoints, mask) ~ detectAndCompute(image, mask, keypoints, noArray(), false)`
 - `compute(image, keypoints, descriptors) ~ detectAndCompute(image, noArray(), keypoints, descriptors, true)`

Step 3: Matching

- **cv::DescriptorMatcher**- abstract class
- Object recognition:
 - Compile keypoints for a variety of objects, save in a database, called a *dictionary* → this stage is called *training*
 - For each new image (*query* image), extract the keypoints
 - Compare these keypoints to the dictionary to find the best match
 - Compare one single descriptor list to a set of descriptor lists (dictionary)
- Tracking:
 - Find all points in this frame of video
 - Find matches for keypoints in the previous frame
 - Estimate how much and where to each object/ person/ vehicle has moved
 - Compare two descriptor lists

Matching functions

- A set of descriptors in the form of a `cv::Mat` object
 - N rows (number of descriptors/ keypoints)
 - D columns (dimensionality of each descriptor, different features describing each keypoint)
- Three functions:
 - `match()`: each keypoint on the query list will be matched to the “best match” from the train list
 - `knnMatch()`: find k nearest neighbors, finds the top k matches
 - `radiusMatch()`: find the best matches within a distance

Table 16-2. Available metrics for the brute force matcher, with their associated formulas; the summations are over the dimensions of the feature vector


Metric	Function
NORM_L2	$\text{dist}(\vec{a}, \vec{b}) = \left[\sum_i (a_i - b_i)^2 \right]^{1/2}$
NORM_L2SQR	$\text{dist}(\vec{a}, \vec{b}) = \sum_i (a_i - b_i)^2$
NORM_L1	$\text{dist}(\vec{a}, \vec{b}) = \sum_i \text{abs}(a_i - b_i)$
NORM_HAMMING	$\text{dist}(\vec{a}, \vec{b}) = \sum_i (a_i \neq b_i) ? 1 : 0$
NORM_HAMMING2	$\text{dist}(\vec{a}, \vec{b}) = \sum_{i(\text{even})} [(a_i \neq b_i) \& \& (a_{i+1} \neq b_{i+1})] ? 1 : 0$

Template Matching

Template Matching [3]

- Searching within an image for a match with a template (often a smaller image)
- Applications:
 - Object recognition
 - Face identification
 - Sign detection
 - Robotics
 - Tracking
 - Augmented reality

Example



Example



0123456789

Template Matching algorithm

- For every possible position (i, j) of the object in the image
 - Evaluate a matching metric and store it in a matching space: $M(i, j)$
 - Search in the matching space for
 - A position with $M(i, j) > T$ (a threshold)
 - Where also $M(i, j)$ is a local maxima (more than neighboring values)
- If using a distance measure, look for minima

0	0	0	0	0	1	0	0
0	0	0	0	0	1	0	0
0	0	0	1	1	1	1	0
0	0	0	1	1	0	1	0
0	0	0	1	0	0	1	0
0	0	0	1	0	0	1	0
0	0	0	1	1	1	1	0
0	0	0	0	0	0	0	0

Image

1	1	1	1
1	0	0	1
1	0	0	1
1	0	0	1
1	1	1	1

Template

11	13	14	11	14
10	14	16	8	14
9	12	12	1	10
10	14	15	7	15

Matching Space

[3] A binary image and a binary template compared using the sum of the squared differences. The best match between the template and the image is towards the bottom right of the matching space where a matching score of 1 was obtained. Note that the comparison is only made where the template fits completely within the image and hence the matching space is smaller than the image.

Matching metrics

- Minimize total differences

$$D_{\text{SquareDifferences}}(i, j) = \sum_{(m, n)} (f(i + m, j + n) - t(m, n))^2$$

$$D_{\text{NormalisedSquareDifferences}}(i, j) = \frac{\sum_{(m, n)} (f(i + m, j + n) - t(m, n))^2}{\sqrt{\sum_{(m, n)} f(i + m, j + n)^2 \sum_{(m, n)} t(m, n)^2}}$$

- Maximize similarity

$$D_{\text{CrossCorrelation}}(i, j) = \sum_{(m, n)} f(i + m, j + n) \cdot t(m, n)$$

$$D_{\text{NormalisedCrossCorrelation}}(i, j) = \frac{\sum_{(m, n)} f(i + m, j + n) \cdot t(m, n)}{\sqrt{\sum_{(m, n)} f(i + m, j + n)^2 \sum_{(m, n)} t(m, n)^2}}$$

*In OpenCV, template matching is supported using normalised cross-correlation **CV_TM_CCORR_NORMED** (although other measures are also supported, such as cross-correlation **CV_TM_CCORR**, sum of the squared differences **CV_TM_SQDIFF** and the normalised sum of the squared differences **CV_TM_SQDIFF_NORMED**):*

```
Mat matching_space;  
matching_space.create(  
    search_image.cols-template_image.cols+1,  
    search_image.rows-template_image.rows+1,  
    CV_32FC1);  
matchTemplate (search_image, template_image,  
               matching_space,  
               CV_TM_CCORR_NORMED);
```


Variations in the images

- When matching:
 - For every possible size
 - For every possible rotation
 - For every possible position (i, j)
 - Calculate the matching measure between image and template
 - Very expensive computation!
- Including more than one template for each object
 - For example if the goal is to locate cars in images, use a dictionary consisting of different car models, colors, perspectives, etc.
 - May contain different sizes, rotations


Chamfer Matching

Chamfer matching

- Designed to lower dependence on the appearance of object of interest




image



template

For example, looking for a square in the image; if they are not the exact same size, the template matching method will fail. Chamfer matching, however, will detect it.

- Based on distance, instead of perfect match



Chamfering algorithm

- The chamfer value for a pixel is the distance from that pixel to the nearest object or edge point



*In OpenCV, to compute a chamfer image from a binary edge image we must invert the edge image so that edges are zero points (for the sake of the **distanceTransform** routine):*

```
Canny (gray_image, edge_image, 100, 200, 3);  
threshold (edge_image, edge_image, 127, 255,  
THRESH_BINARY_INV);  
distanceTransform (edge_image, chamfer_image,  
CV_DIST_L2, 3);
```

Distance Transform

- The distance transform of an image is defined as a new image in which every output pixel is set to a value equal to the distance to the nearest zero pixel in the input image.

```
void cv::distanceTransform(  
    cv::InputArray  src,           // Input image  
    cv::OutputArray dst,          // Result image  
    int             distanceType,  // Distance metric to use  
    int             maskSize      // Mask to use (3, 5, or see below)  
);
```

3.8	2.8	2.4	2	1	0	1	2
3.4	2.4	1.4	1	1	0	1	1.4
3	2	1	0	0	0	0	1
3	2	1	0	0	1	0	1
3	2	1	0	1	1	0	1
3	2	1	0	1	1	0	1
3	2	1	0	0	0	0	1
3.4	2.4	1.4	1	1	1	1	1.4

Chamfer Image

1	1	1	1
1			1
1			1
1			1
1	1	1	1

Template

27.4	19.6	12.8	8	27.4
23.2	16.8	10.4	5	9.4
21	14	8	0	6
23.2	16.8	10.4	5	9.4

Matching Space

Comparing binary edge images (or any binary images)

Matching score = sum of chamfer values for nonzero pixels of template

Minimum score indicates best match

See [3], page 139 for code

References

- [1] Computer Vision: Algorithms and Applications, R. Szeliski
(<http://szeliski.org/Book>)
- [2] Learning OpenCV 3, A. Kaehler & G. Bradski
 - Available online through Safari Books, Seneca libraries
 - https://senecacollege-primo.hosted.exlibrisgroup.com/primo-explore/fulldisplay?docid=01SENC_ALMA5153244920003226&context=L&vid=01SENC&search_scope=default_scope&tab=default_tab&lang=en_US
- [3] Practical introduction to Computer Vision with OpenCV, Kenneth Dawson-Howe
 - Available through Seneca libraries
 - https://senecacollege-primo.hosted.exlibrisgroup.com/primo-explore/fulldisplay?docid=01SENC_ALMA5142810950003226&context=L&vid=01SENC&search_scope=default_scope&tab=default_tab&lang=en_US